



# Observability In Modern Software Architectures

**Mohammad Ali Tofighi**

[ali.tofighi@tapsi.cab](mailto:ali.tofighi@tapsi.cab)

[linkedin.com/in/alitou](https://linkedin.com/in/alitou)

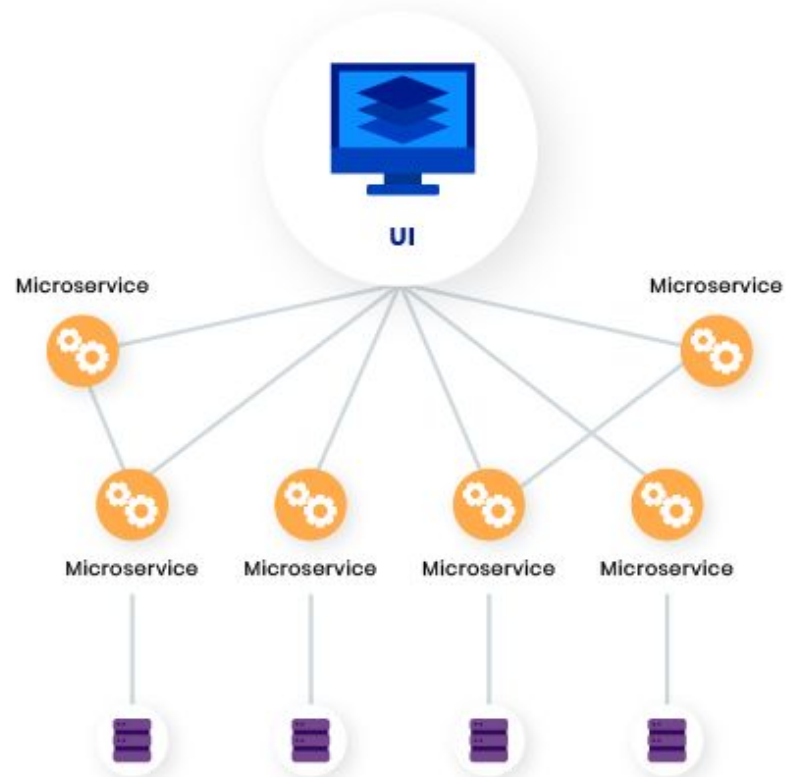
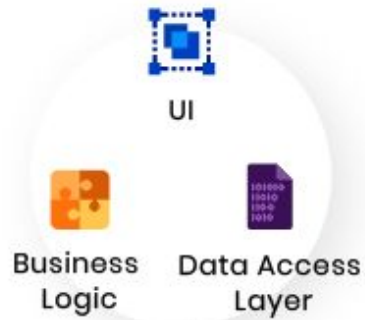
Feb. 2021

- World Gets Complicated
  - Evolution of Software Applications
  - Revolution of Software Architectures
- Three Generations of Application Monitoring
- What Does Observability Mean?
- Three Pillars of Observability
  - Traces
  - Metrics
  - Logs
- Fatal Flaws and Challenges
- Open-Source Tools

# World Gets Complicated

- Simple Applications
  - Observing is trivial
- Complex Applications
  - Different Components
    - Hardware
    - Software
    - Network
    - Load

- Sample Application: Find First N Prime Numbers
  - Starting from simple structures
    - Processing stdin input and print result to stdout
      - The only present factor is our implementation.
  - Getting into more complex ones
    - Getting input from 3rd. party API and store result to database
      - Network overhead and database usage come to play.
    - Calling 1,000 times per second
      - Load affects performance.
- What will you do if you **feel** that your application is **slow**?



- Distribution
  - Introduction of SOA (Service-Oriented Architecture) in 2009
- Abstraction (Cloud)
  - Reduced visibility of application
- Agile Development
  - Resulting in application instability and more living bugs
- Outsourcing
  - Relying on 3rd. party black-box services

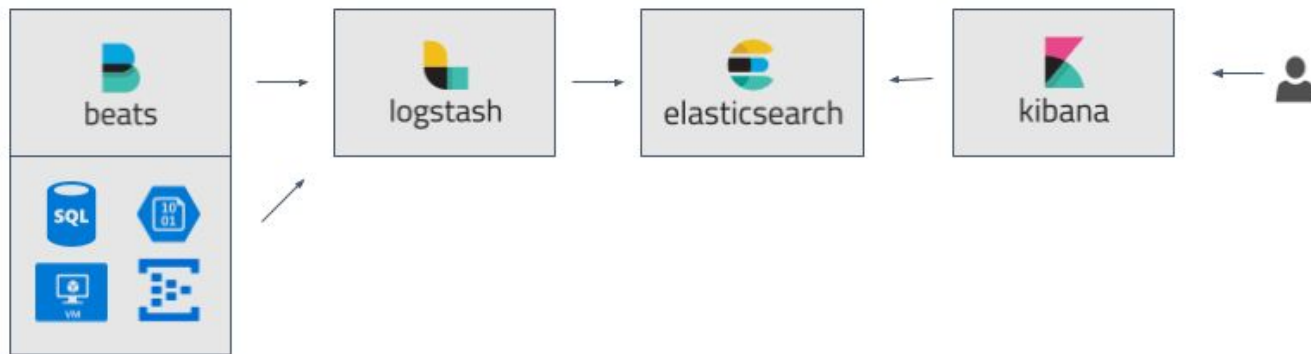
# Three Generations of Application Monitoring



- 1st. Generation – Watching Infrastructure (1990s and early 2000s)
  - Focus on servers uptime, load, network, and storage
  - Dump important logs into text files
  - Analyzing data manually to detect the cause of an issue

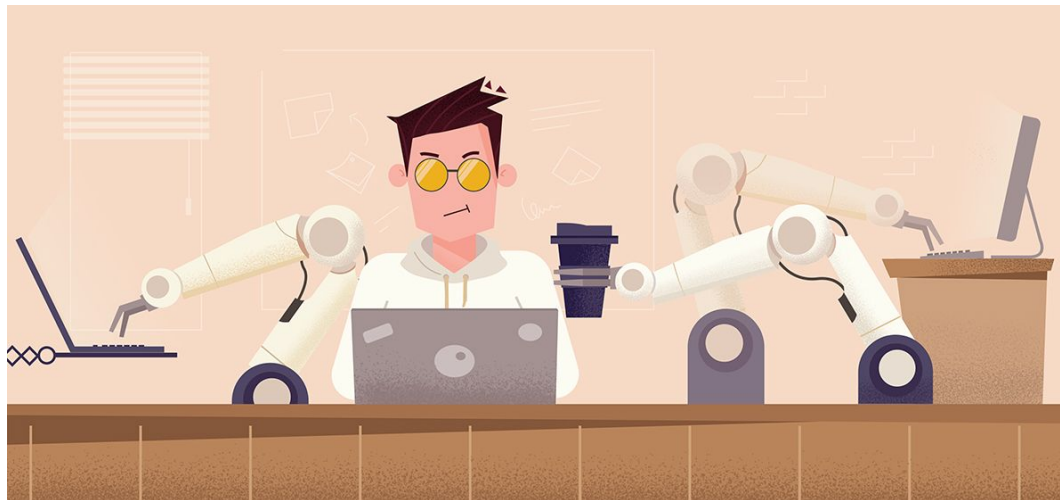
	Time	Level	Thread	Message
	10:48:06....	INFO	Thread_13	Transaction placed successfully.
	10:48:06....	INFO	Thread_3	Client '.Alice' transaction complete.
	10:48:11....	INFO	Thread_11	Verifying network connection.
	10:48:12....	INFO	Thread_7	Client '.Alice' transaction initiated.
	10:48:12....	TRACE	Thread_13	Process start client initiated.
	10:48:12....	TRACE	Thread_6	ESENT database transaction completed.
	10:48:12....	INFO	Thread_1	Network connection established.
	10:48:12....	TRACE	Thread_13	Executing SQL statements.
	10:48:12....	WARN	Thread_10	Transaction field 698 contains invalid contract details.
	10:48:13....	WARN	Thread_8	Transaction time delayed significantly. This may effect available pricing.
	10:48:13....	INFO	Thread_1	Transaction details:
	10:48:14....	ERROR	Thread_2	Invalid transaction details detected. Verifying with server.
	10:48:14....	INFO	Thread_5	Client '.Alice' transaction complete.
	10:48:19....	INFO	Thread_14	Verifying network connection.
	10:48:20....	INFO	Thread_1	Client '.Alice' transaction initiated.

- 2nd. Generation – Gathering Information (2000s and 2010s)
  - Introduction of log management tools, error aggregation tools, notification tools, etc.
  - Getting a lot more information, faster
  - Costing lots of money and time to buy and maintain these tools
  - Still, hard to correlate problems and causes



**What do you guess  
about the 3rd. generation? 🤔**

- 3rd. Generation – Contextual Intelligence
  - Appearance of AIOps in 2016
  - Unify all those tools into an integrated and more ‘aware’ platform
  - Fully automated, from deployment to discovery, problem identification and root cause.
  - NoOps is coming!



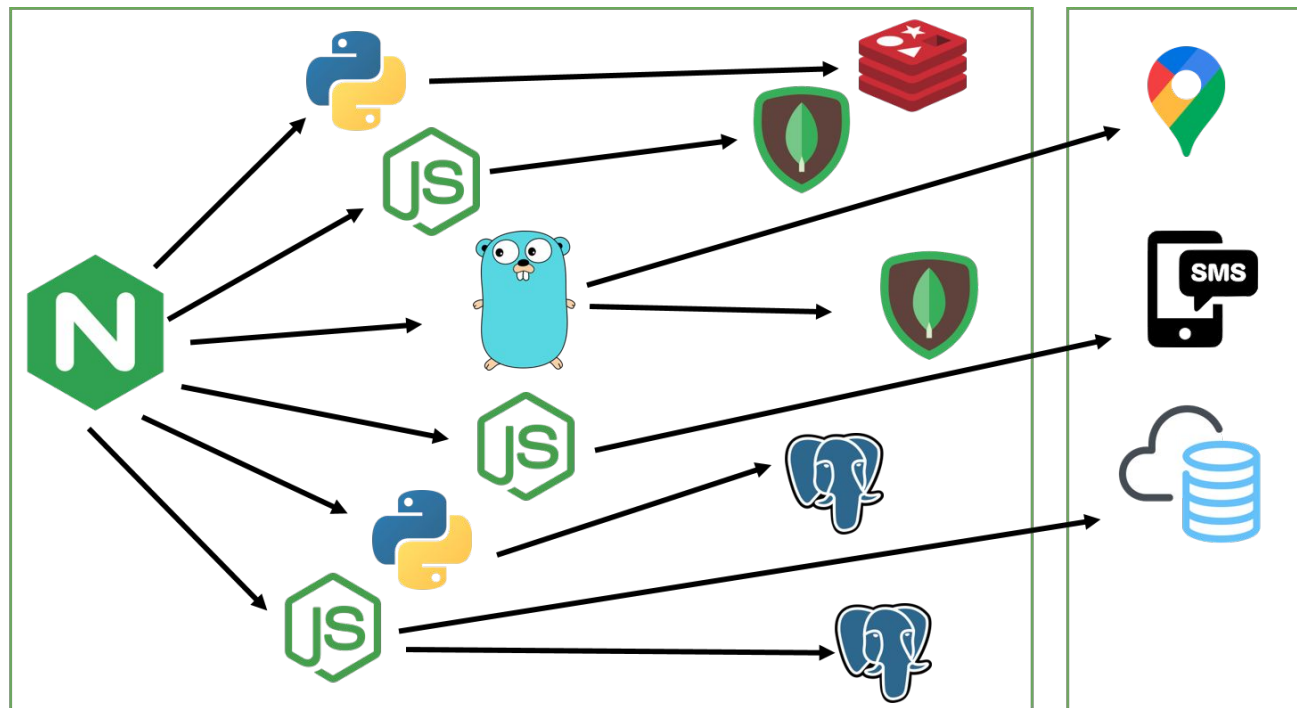
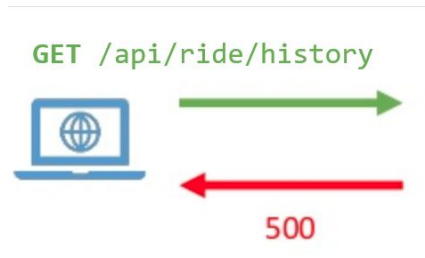
# What Does Observability Mean?

- Originated From Control Theory
  - Observability is a measure of how well **internal states** of a system can be inferred from knowledge of its **external outputs**.



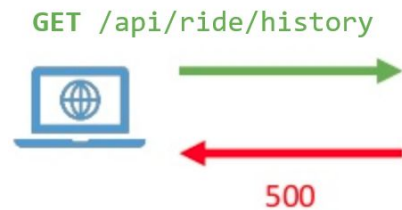
### A sample microservice-based application

What it really is:



A sample microservice-based application

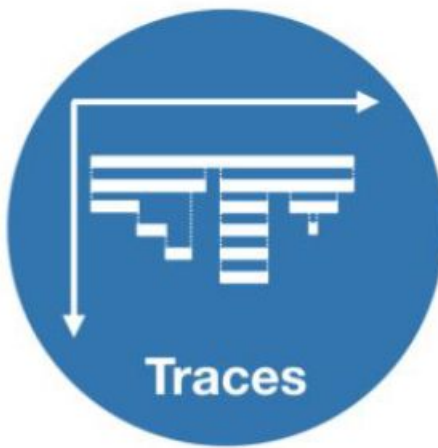
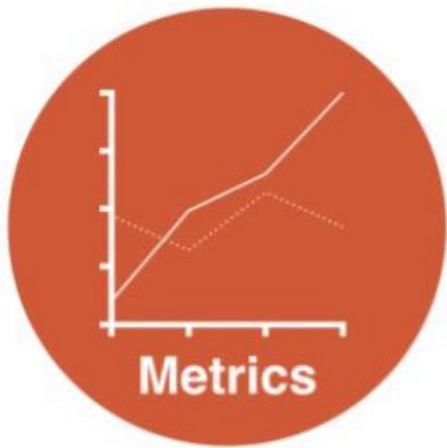
What it feels like:





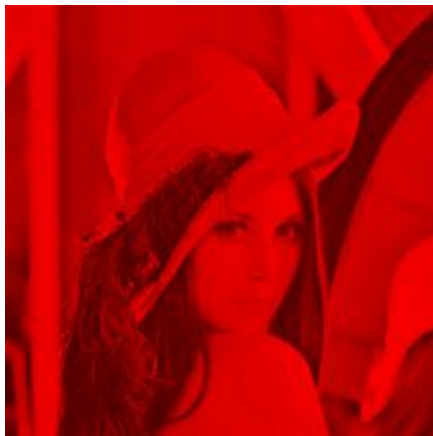
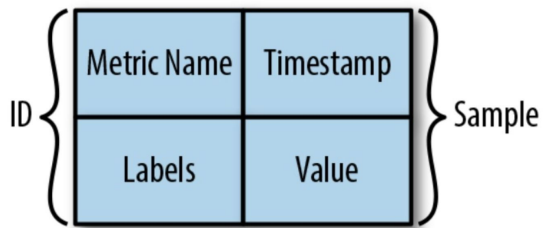
# Three Pillars of Observability

I would like to call them "Three Lenses of Observability"



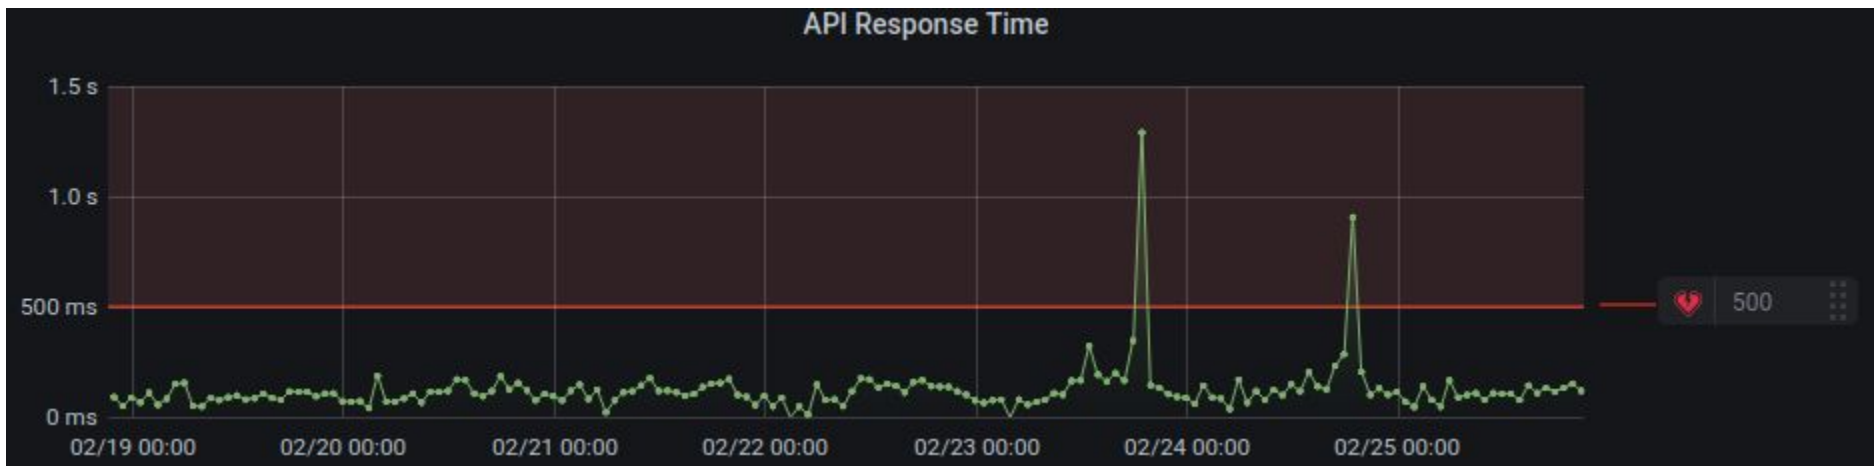
Observability is a measure of how well **internal states** of a system can be inferred from knowledge of its **external outputs**.

- Metrics
  - Reflect the state of the system as a time series of numbers
  - Can be used to trigger an alert when the metric exceeds specified thresholds for a predefined length of time.



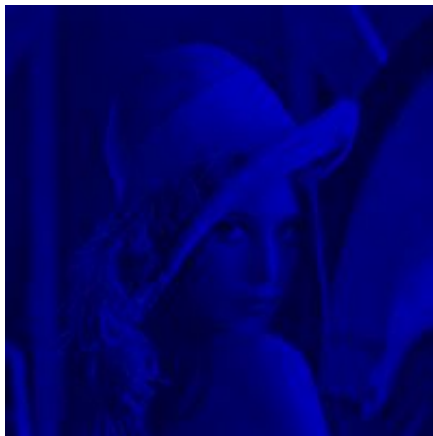
Observability is a measure of how well **internal states** of a system can be inferred from knowledge of its **external outputs**.

- Metrics
  - *Something* happened during the two peaks below



Observability is a measure of how well **internal states** of a system can be inferred from knowledge of its **external outputs**.

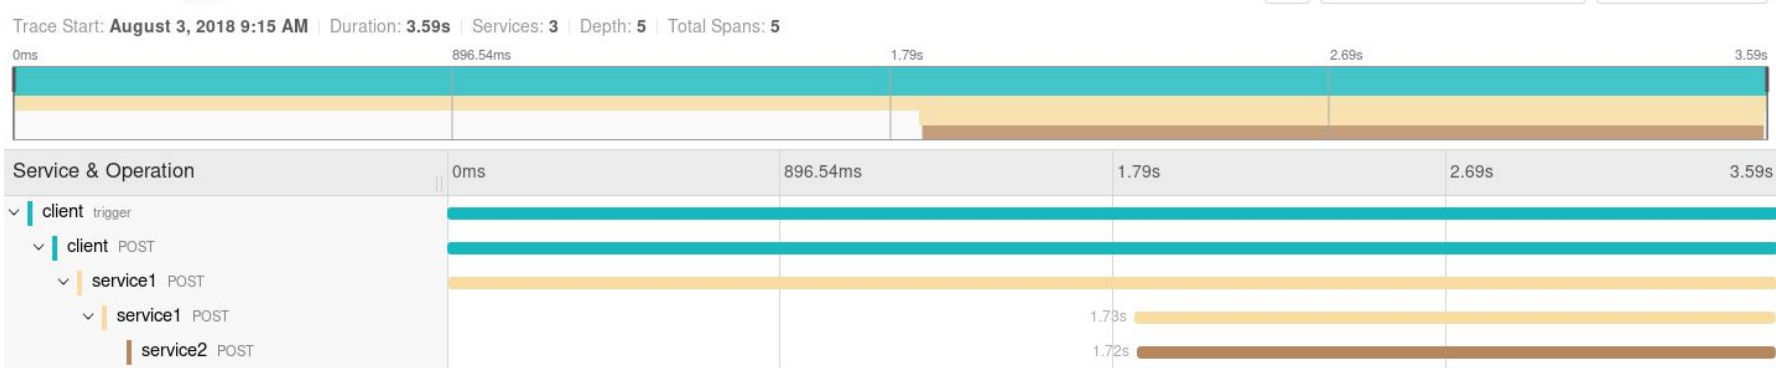
- Traces
  - Show the activity for a transaction within the entire application — all the way from the browser or mobile device down through to the database and back
  - Are the only way to understand the relationships between microservices



Observability is a measure of how well **internal states** of a system can be inferred from knowledge of its **external outputs**.

- Traces
  - *Something slowed down Service 1's work before calling service 2*

### ▼ client: trigger



Observability is a measure of how well **internal states** of a system can be inferred from knowledge of its **external outputs**.

- Logs
  - Help us examine what really happened at the system or software level



Observability is a measure of how well **internal states** of a system can be inferred from knowledge of its **external outputs**.

- Logs
  - A special query in the database didn't use any indexes!

```
message
mongodb 2021-02-12T00:26:01.949+0330 I COMMAND [conn26675] command [redacted] command: find { find: [redacted]
[redacted], projection: {}, returnKey: false, showRecordId: false, lsid: { id: UUID("9cec27de-af2f-4385-a6
e5-5b4225cfd3fb") }, $clusterTime: { clusterTime: Timestamp(1613076960, 59), signature: { hash: BinData(0, BBF15FDC70B9233F77C2FDC878314650913995A6), keyId: 6869354269267460097
}, $db: [redacted] } planSummary: COLLSCAN keysExamined:0 docsExamined:41561 hasSortStage:1 cursorExhausted:1 numYields:324 nreturned:0 reslen:244 locks:{ G
lobal: { acquireCount: { r: 650 } }, Database: { acquireCount: { r: 325 } }, Collection: { acquireCount: { r: 325 } } } protocol:op_msg 115ms
```

Observability is a measure of how well **internal states** of a system can be inferred from knowledge of its **external outputs**.



- Metrics + Traces + Logs: The Greater Sum of the Parts
  - In order to see the full picture, we need to use all the data.



Observability is a measure of how well **internal states** of a system can be inferred from knowledge of its **external outputs**.

# Fatal Flaws and Challenges

- Dependency of "Pillars"
  - Metrics, Traces, and Logs are just pieces, with low benefits on their own.
  - None of them directly addresses a particular pain point, use case, or business need.

- "Metrics are enough. Tracing and log aggregation are hard to implement!"
  - If `service="foo"` has problems, we can check metrics with the same label
  - Usually, there are more than *thousands* of metrics for `service="foo"`
  - Metrics aren't enough!

- "We use log aggregators because devs shouldn't have access to servers!"
  - No. We use them because of:
    - Centralization
    - Searchability
    - Accessibility
    - Monitoring & Alerting

- The Huge Amount of Logs

- If we want to use logs to account for individual transactions:

- Transaction rate ×
    - Number of microservice instances ×
    - Network and storage cost ×
    - Weeks of data retention

- Too much cost!

- Logs are less useful than monolithic applications

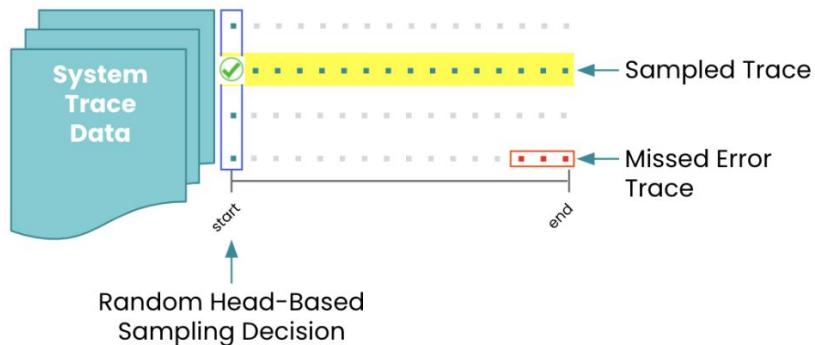
- Due to distributed and concurrent nature of microservices



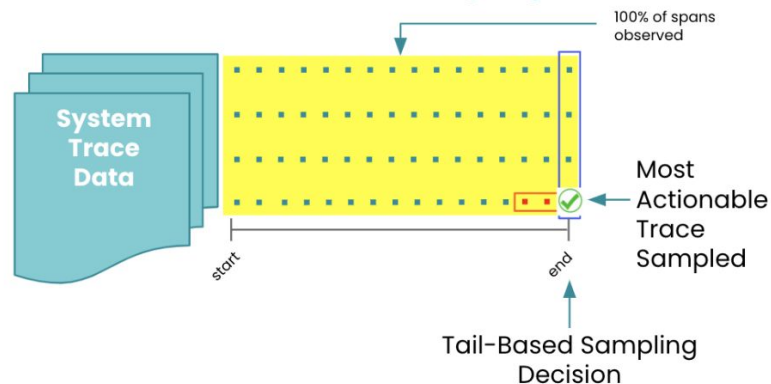
- What To Monitor?
  - Google's Golden Signals: <https://sre.google/sre-book/monitoring-distributed-systems>
  - RED
    - Rate
    - Errors
    - Duration
  - USE
    - Utilization
    - Saturation
    - Errors

- What To Trace? How Often Do The Tracing?
  - Head-based sampling
  - Tail-based sampling

### Traditional Head-Based Sampling



### Tail-Based Sampling





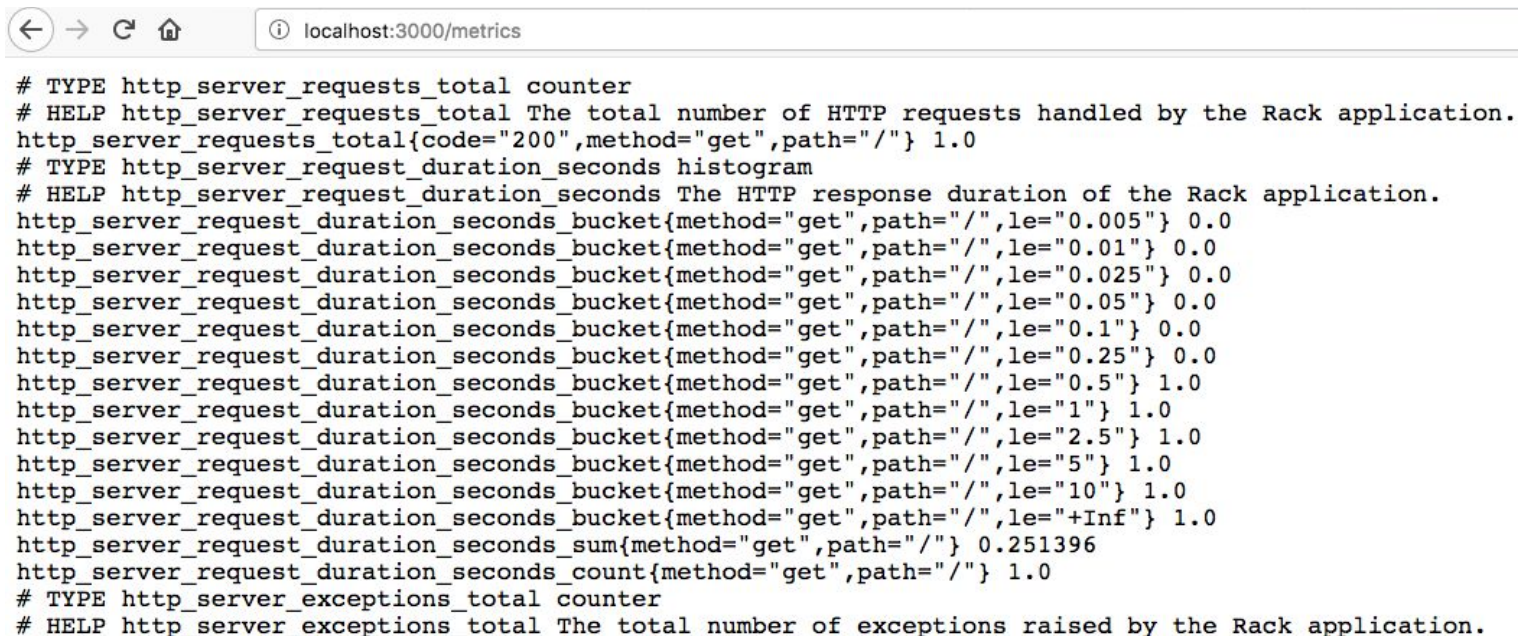
- Yellow = Green + Red
  - Requires experience or specific knowledge
  - Correlating metrics, traces and logs is really an extraordinary complicated task
  - We may make mistakes during correlating them, and confuse ourselves
  - We will need machines. Remember the 3rd. generation of monitoring

# Open-Source Tools

- Prometheus
  - Is an open-source systems monitoring and alerting toolkit originally built at SoundCloud in 2012;
  - Joined the Cloud Native Computing Foundation in 2016;
  - Is the second CNCF graduated project, after Kubernetes;
  - Has a multi-dimensional data model with time series data identified by metric name and key/value pair of labels.



- Prometheus



The screenshot shows a web browser window with the address bar displaying 'localhost:3000/metrics'. The page content is a text-based representation of Prometheus metrics. It includes comments for the metric types (counter and histogram) and their descriptions, followed by the actual metric values for various HTTP server metrics. The metrics are listed in a single column, with each line representing a different metric or a specific bucket within a histogram.

```
# TYPE http_server_requests_total counter
# HELP http_server_requests_total The total number of HTTP requests handled by the Rack application.
http_server_requests_total{code="200",method="get",path="/" } 1.0
# TYPE http_server_request_duration_seconds histogram
# HELP http_server_request_duration_seconds The HTTP response duration of the Rack application.
http_server_request_duration_seconds_bucket{method="get",path="/",le="0.005"} 0.0
http_server_request_duration_seconds_bucket{method="get",path="/",le="0.01"} 0.0
http_server_request_duration_seconds_bucket{method="get",path="/",le="0.025"} 0.0
http_server_request_duration_seconds_bucket{method="get",path="/",le="0.05"} 0.0
http_server_request_duration_seconds_bucket{method="get",path="/",le="0.1"} 0.0
http_server_request_duration_seconds_bucket{method="get",path="/",le="0.25"} 0.0
http_server_request_duration_seconds_bucket{method="get",path="/",le="0.5"} 1.0
http_server_request_duration_seconds_bucket{method="get",path="/",le="1"} 1.0
http_server_request_duration_seconds_bucket{method="get",path="/",le="2.5"} 1.0
http_server_request_duration_seconds_bucket{method="get",path="/",le="5"} 1.0
http_server_request_duration_seconds_bucket{method="get",path="/",le="10"} 1.0
http_server_request_duration_seconds_bucket{method="get",path="/",le="+Inf"} 1.0
http_server_request_duration_seconds_sum{method="get",path="/" } 0.251396
http_server_request_duration_seconds_count{method="get",path="/" } 1.0
# TYPE http_server_exceptions_total counter
# HELP http_server_exceptions_total The total number of exceptions raised by the Rack application.
```

- Grafana
  - Is a multi-platform open source analytics and interactive visualization tool, started in 2014;
  - Allowing to query, visualize, alert on and understand metrics, no matter where they are stored.



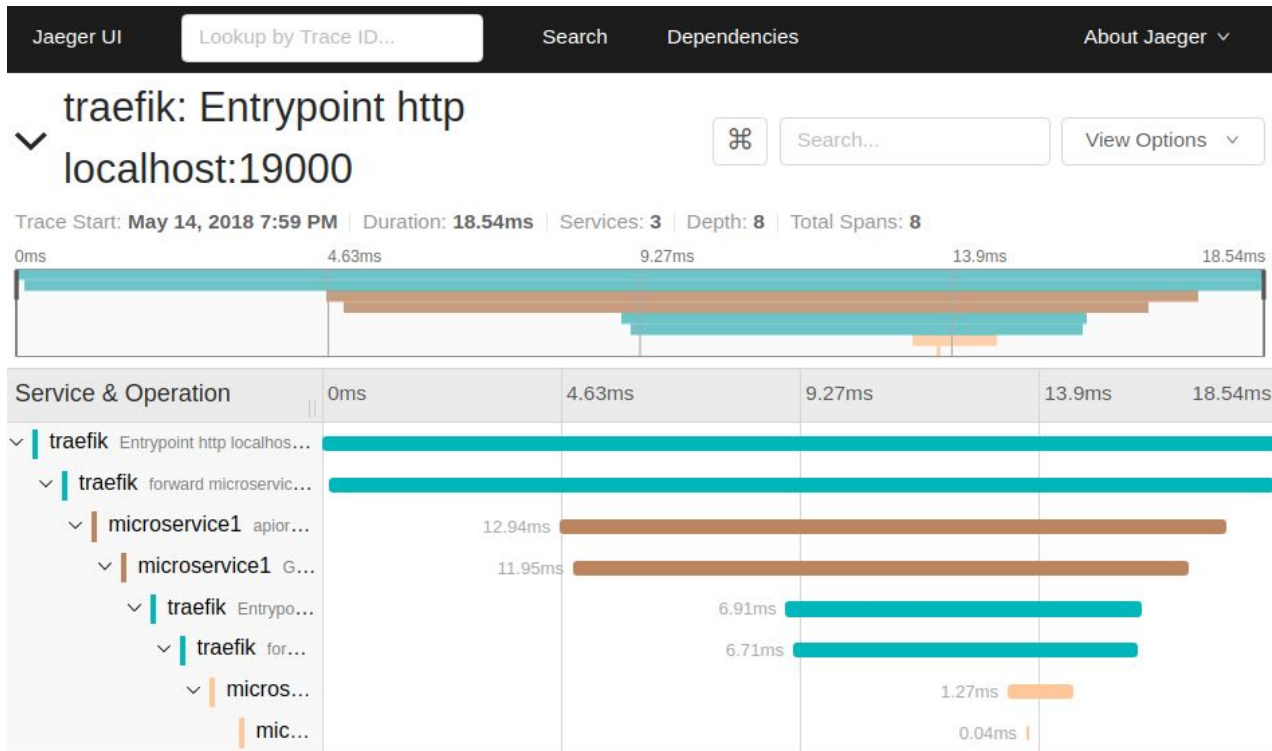
- Grafana



- Jaeger
  - Is an open source, end-to-end distributed tracing tool initially developed at Uber in 2017;
  - Is the 7th. graduated projects from CNCF.



- Jaeger



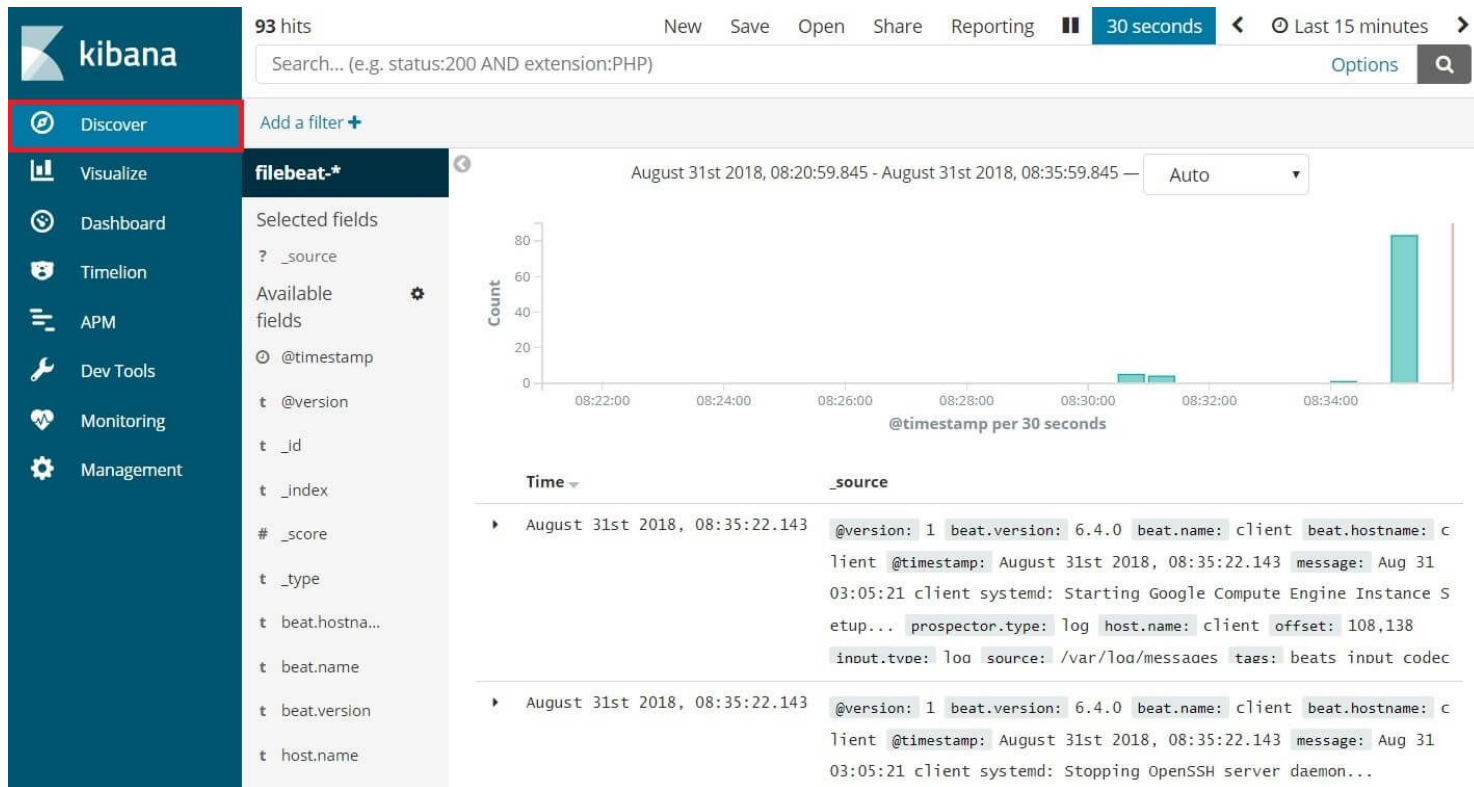


- ELK Stack

- "ELK" is the acronym for three open source projects: Elasticsearch, Logstash, and Kibana;
- Elasticsearch is a search and analytics engine;
- Logstash is a data processing pipeline that ingests data from multiple sources and then sends it to Elasticsearch;
- Kibana lets users visualize data with charts and graphs in Elasticsearch.



- ELK Stack



- There are more!

graylog



fluentbit

ZABBIX



ZIPKIN



fluentd

Nagios®



# Thanks!

[ali.tofighi@tapsi.cab](mailto:ali.tofighi@tapsi.cab)